



MANIPAL INSTITUTE OF TECHNOLOGY
MANIPAL
(A constituent unit of MAHE, Manipal)

PCAP: Mini Project Report on

Parallel Encryption, Cracking and Decryption: A CUDA Approach to the Vigenère Cipher

SUBMITTED BY

SUJEET SANJAY AMBERKAR 200905092

Under the Guidance of:

Mrs Neelima Bayyapu

**Department of Computer Science and Engineering
Manipal Institute of Technology, Manipal Karnataka – 576104**

April 2024

Index

Sr NO	Heading	Page Number
1	Abstract	3
2	Introduction	3
3	Project Objective	3
4	Understanding CUDA	4
5	Selecting Algorithm	4
6	Normal C Code	5-11
7	CUDA Code	12-19
8	Implementing CUDA Kernels	20
9	Optimization Strategies	21
10	Memory Management	21
11	Execution Time Measurements	21
12	Bottlenecks and Limitations	22
13	References	22
14	Conclusion	22

Abstract:

This project explores the encryption cracking and decryption of Vigenère cipher using the power of parallel computing. This study leverages the Compute Unified Device Architecture (CUDA) to implement and optimize the encryption, decryption, and cryptanalysis of the Vigenère cipher on NVIDIA GPUs. The results show a significant increase in the speed of cryptanalysis when done using CUDA.

Introduction:

In the field of cryptography, the ability to process large volumes of data with high efficiency and speed is very important. CUDA's architecture is designed for Single Instruction Multiple Data. Cryptographic algorithms for Encryption, Decryption and cracking of ciphers often involve similar or identical operations across all data points making it an ideal fit for using CUDA to increase the speed. The Shared memory also helps us to resources between threads thus reducing the latency and increasing the bandwidth. Importantly, in cryptographic systems such as the Vigenère cipher, the size of the key space expands exponentially with the length of the key, following a $26^{\text{(size of the key)}}$ pattern. This immense key space underscores the need for CUDA's powerful parallel processing capabilities, as it enables the exploration of vast cryptographic landscapes within feasible timeframes,

Project Objective

The Vigenère cipher is a method of encrypting alphabetic text where each letter of the plaintext is encoded with a different Caesar cipher. It is an evolution of the Caesar cipher in which we conduct a series of different shifts or keys rather than just a single key in the Caesar cipher.

Core Objectives:

- 1) **CUDA- Enabled Vigenère Encryption and Decryption:** This involves developing a framework using CUDA for performing Encryption and Decryption taking advantage of GPU Architecture and Shared Memory in CUDA.
- 2) **Cryptanalysis Techniques:** Implement advanced CUDA-based implementation of the Kasiski algorithm to guess the key length and then find the key of Encrypted text using Frequency Analysis by parallel computing.
- 3) **Comparative Performance Analysis:** Conduct a thorough performance analysis comparing the CUDA-accelerated Vigenère cipher operations and cryptanalysis methods against traditional CPU-based implementations. This comparative study aims to highlight the efficiency and speed gains achieved through GPU parallelism,

Understanding CUDA

- 1) **Parallel Execution with CUDA Kernels:** The CUDA kernel such as "repeatString", "vignere_encrypt_cuda", "crack_vigenere_kernel", "vignere_decrypt_cuda", "calculate_matches" leverages the power of Single Instruction Multiple Data by performing similar instructions across across multiple data points
- 2) **Efficient Memory Usage:** The use of `__constant__` memory types, as seen with `d_inputString` and `d_key`, optimizes memory. **cudaMemcpyToSymbol** is specifically used to copy the encryption key and decryption key to these memory spaces. Ensuring all threads have quick access to these constants.
- 3) **Optimal Resource Utilization:** The Effective Kernel launch configurations (`<<<numBlocks, blockSize>>>`) for the CUDA process.

Selecting Algorithm

1) Encryption:

$$E_i(x_i) = (x_i + k_i) \bmod 26$$

Each character in the plaintext is shifted according to the corresponding character in the key. This operation, due to its simplicity and repetitive nature across the length of the plaintext, is ideally suited for GPU's SIMD. Allowing encryption of large texts in parallel with significant speed improvements

2) Decryption:

$$E_i(x_i) = (x_i - k_i) \bmod 26$$

This is exactly inverse of Encryption. Similar to encryption, decryption can be efficiently parallelized over the GPU, with each thread handling the decryption of individual characters or blocks of text, thus benefiting from CUDA's parallel processing capabilities.

3) **Key Length Determination (Kasiski Examination):** This method involves analyzing the ciphertext to find repeating sequences of characters and measuring the distances between these repetitions. Since we shift and repeat the process again, CUDA can significantly reduce the time required by performing these analyses in parallel across multiple segments of the text.

4) **Frequency Analysis:** Once the key length is known, the cipher can be treated as multiple Caesar ciphers. Frequency analysis is then applied to each segment in parallel to get the key.

The chosen algorithms are inherently parallelizable, making them excellent candidates for CUDA's architecture

Normal C Code:

```
#include <stdio.h>
#include <string.h>
#include <stdlib.h>
#include <ctype.h>
#include <math.h>
#include <time.h>
#define NUM_LETTERS 26
char *only_alphabets(const char *text)
{
    char *res = malloc(strlen(text) + 1);
    if (res == NULL)
    {
        printf("Memory allocation failed.\n");
        exit(1);
    }

    int j = 0;
    for (int i = 0; text[i] != '\0'; i++)
    {
        if (isalpha((unsigned char)text[i]))
        {
            res[j++] = text[i];
        }
    }
    res[j] = '\0';

    return res;
}
typedef struct
{
    int value;
    int index;
} ElementIndexPair;
```

```

ElementIndexPair *zip_index(int *arr, int size, int *outSize)
{
    ElementIndexPair *result = (ElementIndexPair *)malloc(size *
sizeof(ElementIndexPair));
    if (result == NULL)
    {
        printf("Memory allocation failed.\n");
        exit(1);
    }

    for (int i = 0; i < size; i++)
    {
        result[i].value = arr[i];
        result[i].index = i;
    }

    *outSize = size;

    return result;
}

void get_char_freq(const char *text, int freq[26])
{
    for (int i = 0; i < 26; i++)
    {
        freq[i] = 0;
    }

    while (*text)
    {
        if (isalpha((unsigned char)*text))
        {
            int index = toupper((unsigned char)*text) - 'A';
            freq[index]++;
        }
        text++;
    }
}

```

```

void vignere_encrypt(const char *msg, const char *key, char *encrypted)
{
    int msgLen = strlen(msg);
    int keyLen = strlen(key);

    for (int i = 0; i < msgLen; i++)
    {
        if (isupper(msg[i]))
        {
            int msgIndex = msg[i] - 'A';
            int keyIndex = key[i % keyLen] - 'A';
            int encryptedIndex = (msgIndex + keyIndex) % NUM_LETTERS;
            encrypted[i] = 'A' + encryptedIndex;
        }
        else
        {
            encrypted[i] = msg[i];
        }
    }
    encrypted[msgLen] = '\0';
}

void vignere_decrypt(const char *encrypted_msg, const char *key, char *decrypted)
{
    int msgLen = strlen(encrypted_msg);
    int keyLen = strlen(key);

    for (int i = 0; i < msgLen; i++)
    {
        if (isupper(encrypted_msg[i]))
        {
            int msgIndex = encrypted_msg[i] - 'A';
            int keyIndex = key[i % keyLen] - 'A';
            int decryptedIndex = (msgIndex - keyIndex + NUM_LETTERS) %
NUM_LETTERS;
            decrypted[i] = 'A' + decryptedIndex;
        }
    }
}

```

```

else
{
    decrypted[i] = encrypted_msg[i];
}

}

decrypted[msgLen] = '\0';
}

double mean(const int arr[], int size)
{
    double sum = 0;
    for (int i = 0; i < size; i++)
    {
        sum += arr[i];
    }
    return sum / size;
}

double variance(const int arr[], int size, double mean)
{
    double sum = 0;
    for (int i = 0; i < size; i++)
    {
        sum += (arr[i] - mean) * (arr[i] - mean);
    }
    return sum / size;
}

int mode(const int arr[], int size)
{
    int maxValue = 0, maxCount = 0;

    for (int i = 0; i < size; ++i)
    {
        int count = 0;
        for (int j = 0; j < size; ++j)
        {
            if (arr[j] == arr[i])
                ++count;
        }
    }
}

```



```

        if (count > maxCount)
        {
            maxCount = count;
            maxValue = arr[i];
        }
    }
    return maxValue;
}

int guess_key_length(const char *cipher)
{
    int length = strlen(cipher);
    int *res = (int *)malloc(length * sizeof(int));
    for (int shift = 0; shift < length; shift++)
    {
        int matches = 0;
        for (int i = 0; i < length; i++)
        {
            if (cipher[i] == cipher[(i - shift + length) % length])
            {
                matches++;
            }
        }
        res[shift] = matches;
    }
    double dat_means = mean(res + 1, length - 1);
    double dat_std_dev = sqrt(variance(res + 1, length - 1, dat_means));
    int *peaks = (int *)malloc(length * sizeof(int));
    int peaks_count = 0;
    for (int i = 0; i < length; i++)
    {
        if (res[i] >= dat_std_dev + dat_means)
        {
            peaks[peaks_count++] = i;
        }
    }
    int *peak_diff = (int *)malloc((peaks_count - 1) * sizeof(int));

```

```

    for (int i = 0; i < peaks_count - 1; i++)
    {
        peak_diff[i] = peaks[i + 1] - peaks[i];
    }

    int key_length = mode(peak_diff, peaks_count - 1);

    free(res);
    free(peaks);
    free(peak_diff);
    return key_length;
}

char *crack_vigenere(const char *cipher)
{
    int n_key = guess_key_length(cipher);
    int cipher_len = strlen(cipher);
    char *keys = (char *)malloc(n_key + 1);

    for (int i = 0; i < n_key; i++)
    {
        int freq[NUM_LETTERS] = {0};
        for (int j = i; j < cipher_len; j += n_key)
        {
            if (isupper(cipher[j]))
            {
                freq[cipher[j] - 'A']++;
            }
        }

        int max_freq = 0;
        char max_char = 0;
        for (int k = 0; k < NUM_LETTERS; k++)
        {
            if (freq[k] > max_freq)
            {
                max_freq = freq[k];
                max_char = k + 'A';
            }
        }
    }
}

```

```

        keys[i] = ((max_char - 'E' + NUM_LETTERS) % NUM_LETTERS) + 'A';
    }
    keys[n_key] = '\0';

    return keys;
}

int main()
{
    clock_t start = clock();

    char *message = "Standard deviation.... underlying risk";
    size_t messageLen = strlen(message);
    char *repeatedMessage = (char *)malloc(messageLen * 10 + 1);
    for (int i = 0; i < 10; i++)
    {
        strcat(repeatedMessage, message);
    }
    char *filteredMessage = only_alphabets(repeatedMessage);
    for (int i = 0; filteredMessage[i]; i++)
    {
        filteredMessage[i] = toupper(filteredMessage[i]);
    }
    const char *key = "SUJ";
    char *encrypted = (char *)malloc(strlen(filteredMessage) + 1);

    vigenere_encrypt(filteredMessage, key, encrypted);
    char *guessedKey = crack_vigenere(encrypted);
    printf("%s\n", guessedKey);
    free(repeatedMessage);
    free(filteredMessage);
    free(encrypted);
    free(guessedKey);
    clock_t end = clock();
    double time_spent = (double)(end - start) / CLOCKS_PER_SEC;
    printf("Time spent: %f seconds\n", time_spent);
    return 0;
}

```

CUDA CODE:

```
#include <cuda_runtime.h>
#include <stdlib.h>
#include <string.h>
#include <ctype.h>
#include <time.h>

#define MAX_LENGTH 10000
#define REPEAT_TIMES 10
#define MAX_KEY 10
#define NUM_LETTERS 26

__constant__ char d_inputString[MAX_LENGTH];
__device__ char d_key[10];

__device__ char d_key_calculated[1024];
__device__ bool is_upper(char c) {
    return c >= 'A' && c <= 'Z';
}

__global__ void repeatString(char *output, int inputLength){
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    int offset = idx * inputLength;
    for(int i = 0; i<inputLength;i++)
        output[offset+i]=d_inputString[i];
}

__global__ void vigenere_encrypt_cuda(const char* msg, char* encrypted, int
msgLen, int keyLen) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < msgLen) {
        if (is_upper(msg[idx])) {
            int msgIndex = msg[idx] - 'A';
            int keyIndex = d_key[idx % keyLen] - 'A';
            int encryptedIndex = (msgIndex + keyIndex) % NUM_LETTERS;
            encrypted[idx] = 'A' + encryptedIndex;
        } else {
            encrypted[idx] = msg[idx];
        }
    }
}
```

```

__global__ void crack_vigenere_kernel(const char* cipher, int cipher_len, int n,
char* keys) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < n) {
        int freq[NUM_LETTERS] = {0};
        for (int j = i; j < cipher_len; j += n) {
            if (is_upper(cipher[j])) {
                freq[cipher[j] - 'A']++;
            }
        }
        int max_freq = 0;
        char max_char = 0;
        for (int k = 0; k < NUM_LETTERS; k++) {
            if (freq[k] > max_freq) {
                max_freq = freq[k];
                max_char = k + 'A';
            }
        }
        keys[i] = ((max_char - 'E' + NUM_LETTERS) % NUM_LETTERS) + 'A';
    }
}

__global__ void vigenere_decrypt_cuda(const char* encrypted, char* decrypted, int
msgLen, int keyLen) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < msgLen) {
        if (is_upper(encrypted[idx])) {
            int encryptedIndex = encrypted[idx] - 'A';
            int keyIndex = d_key[idx % keyLen] - 'A';
            int msgIndex = (encryptedIndex - keyIndex + NUM_LETTERS) %
NUM_LETTERS;
            decrypted[idx] = 'A' + msgIndex;
        } else {
            decrypted[idx] = encrypted[idx];
        }
    }
}

```

```

char* only_alphabets(const char* text) {
    char* res = (char*)malloc(strlen(text) + 1);
    if (res == NULL) {
        printf("Memory allocation failed.\n");
        exit(1);
    }
    int j = 0;
    for (int i = 0; text[i] != '\0'; i++) {
        if (isalpha((unsigned char)text[i]))
        {
            res[j++] = toupper((unsigned char)text[i]);
        }
    }
    res[j] = '\0';

    return res;
}

char* process_string_CUDA(char* temp) {
    temp = only_alphabets(temp);
    char *d_output;
    int inputLength = strlen(temp);

    size_t outputSize = MAX_LENGTH * REPEAT_TIMES;
    cudaMemcpyToSymbol(d_inputString, temp, inputLength + 1);
    cudaDeviceSynchronize();

    cudaMalloc(&d_output, outputSize);
    cudaMemset(d_output, 0, outputSize);
    repeatString<<<REPEAT_TIMES, 1>>>(d_output, inputLength);
    cudaDeviceSynchronize();

    char *plainTextInput = (char*) malloc(outputSize);
    cudaMemcpy(plainTextInput, d_output, outputSize, cudaMemcpyDeviceToHost);
    cudaFree(d_output);
    free(temp);
    return plainTextInput;
}

```

```

char *cudaEncrypt(char *processedInput, char *key)
{
    char* d_message;
    char* d_encrypted;
    char* d_decrypted;
    int msgLen=strlen(processedInput);
    int keyLen = strlen(key);
    cudaMalloc((void**)&d_message, msgLen+1);
    cudaMalloc((void**)&d_encrypted, msgLen+1);
    cudaMalloc((void**)&d_decrypted, msgLen+1);

    cudaMemcpy(d_message, processedInput, msgLen+1, cudaMemcpyHostToDevice);
    cudaMemcpyToSymbol(d_key, key, keyLen+1);

    int blockSize = 256;
    int numBlocks = (msgLen + blockSize - 1) / blockSize;

    vignere_encrypt_cuda<<<numBlocks, blockSize>>>(d_message, d_encrypted,
msgLen, keyLen);
    char* encrypted = (char*)malloc(msgLen+1);
    cudaMemcpy(encrypted, d_encrypted, msgLen+1, cudaMemcpyDeviceToHost);
    cudaFree(d_message);
    cudaFree(d_encrypted);
    cudaFree(d_decrypted);
    return encrypted;
}

__global__ void calculate_matches(const char *cipher, int length, int *matches) {
    int shift = blockIdx.x * blockDim.x + threadIdx.x;
    if (shift >= length) return;

    int matchCount = 0;
    for (int i = 0; i < length; ++i) {
        if (cipher[i] == cipher[(i - shift + length) % length]) {
            ++matchCount;
        }
    }
    matches[shift] = matchCount;
}

```

```

double mean(const int arr[], int size) {
    double sum = 0;
    for (int i = 0; i < size; i++) {
        sum += arr[i];
    }
    return sum / size;
}

double variance(const int arr[], int size, double mean) {
    double sum = 0;
    for (int i = 0; i < size; i++) {
        sum += (arr[i] - mean) * (arr[i] - mean);
    }
    return sum / size;
}

int mode(const int arr[], int size) {
    int maxValue = 0, maxCount = 0;
    for (int i = 0; i < size; ++i) {
        int count = 0;
        for (int j = 0; j < size; ++j) {
            if (arr[j] == arr[i]) ++count;
        }
        if (count > maxCount) {
            maxCount = count;
            maxValue = arr[i];
        }
    }
    return maxValue;
}

int guess_key_length_cuda(const char *cipher) {
    int length = strlen(cipher);
    char *d_cipher;
    int *d_matches, *matches;

    cudaMalloc((void **)&d_cipher, length);
    cudaMalloc((void **)&d_matches, length * sizeof(int));
    cudaMemcpy(d_cipher, cipher, length, cudaMemcpyHostToDevice);

```



```

int threadsPerBlock = 256;
int blocks = (length + threadsPerBlock - 1) / threadsPerBlock;

calculate_matches<<<blocks, threadsPerBlock>>>(d_cipher, length, d_matches);

matches = (int *)malloc(length * sizeof(int));
cudaMemcpy(matches, d_matches, length * sizeof(int), cudaMemcpyDeviceToHost);
double dat_means = mean(matches + 1, length - 1);
double dat_std_dev = sqrt(variance(matches + 1, length - 1, dat_means));

int* peaks = (int*)malloc(length * sizeof(int));
int peaks_count = 0;
for (int i = 0; i < length; i++) {
    if (matches[i] >= dat_std_dev + dat_means) {
        peaks[peaks_count++] = i;
    }
}

int* peak_diff = (int*)malloc((peaks_count - 1) * sizeof(int));
for (int i = 0; i < peaks_count - 1; i++) {
    peak_diff[i] = peaks[i + 1] - peaks[i];
}

int key_length = mode(peak_diff, peaks_count - 1);
cudaFree(d_cipher);
cudaFree(d_matches);
free(matches);

free(peaks);
free(peak_diff);

return key_length;
}

char* crack_vigenere_cuda(const char* cipher, int n) {
    int cipher_len = strlen(cipher);
    char* dev_cipher = NULL;
    char* dev_keys = NULL;
    char* keys = (char*)malloc(n + 1);

```

```

    cudaMalloc((void**)&dev_cipher, cipher_len * sizeof(char));
    cudaMalloc((void**)&dev_keys, n * sizeof(char));

    cudaMemcpy(dev_cipher, cipher, cipher_len * sizeof(char),
cudaMemcpyHostToDevice);

    int threadsPerBlock = 256;
    int blocksPerGrid = (n + threadsPerBlock - 1) / threadsPerBlock;
    crack_vigenere_kernel<<<blocksPerGrid, threadsPerBlock>>>(dev_cipher,
cipher_len, n, dev_keys);

    cudaMemcpy(keys, dev_keys, n * sizeof(char), cudaMemcpyDeviceToHost);
    keys[n] = '\0';

    cudaFree(dev_cipher);
    cudaFree(dev_keys);

    return keys;
}
char* cudaDecrypt(char *encrypted, char *key) {
    char* d_encrypted;
    char* d_decrypted;
    int msgLen = strlen(encrypted);
    int keyLen = strlen(key);

    cudaMalloc((void**)&d_encrypted, msgLen + 1);
    cudaMalloc((void**)&d_decrypted, msgLen + 1);

    cudaMemcpy(d_encrypted, encrypted, msgLen + 1, cudaMemcpyHostToDevice);

    cudaMemcpyToSymbol(d_key_calculated, key, keyLen + 1);

    int blockSize = 256;
    int numBlocks = (msgLen + blockSize - 1) / blockSize;

    vigenere_decrypt_cuda<<<numBlocks, blockSize>>>(d_encrypted, d_decrypted,
msgLen, keyLen);

```

```

    cudaDeviceSynchronize();

    char* decrypted = (char*)malloc(msgLen + 1);

    cudaMemcpy(decrypted, d_decrypted, msgLen + 1, cudaMemcpyDeviceToHost);

    cudaFree(d_encrypted);
    cudaFree(d_decrypted);
    return decrypted;
}

int main ()
{
    clock_t start = clock();
    char inputString[MAX_LENGTH]="Standard deviation ... underlying risk";
    char key[MAX_KEY]="SUJ";
    char * processedInput = process_string_CUDA(inputString);
    char * encrypted = cudaEncrypt(processedInput,key);
    int n = guess_key_length_cuda(encrypted);
    char* key_calculated = crack_vigenere_cuda(encrypted,n);
    char *calculatedPlainText=cudaDecrypt(encrypted,key_calculated);
    printf("%s",key_calculated);
    clock_t end = clock();
    double time_spent = (double)(end - start) / CLOCKS_PER_SEC;
    printf("Time spent: %f seconds\n", time_spent);

    return 0;
}

```

Implementing CUDA Kernels:

1) Encryption Kernel(vignere_encrypt_cuda):

This kernel encrypts input text using the Vigenère cipher. Each thread encrypts a single character of the input string, calculating the corresponding encrypted character by shifting it according to the key. The kernel is designed to handle each character independently, enabling the encryption process to scale efficiently with the length of the text.

2)Decryption Kernel(vignere_decrypt_cuda): The Decryption process is exactly inverse of the encryption process. Thus this follows the same kernel design as the Encryption

3)Finding Length (calculate_matches):

The CUDA kernel `calculate_matches` and the function `guess_key_length_cuda` are central to the implementation of an efficient approach to determining the key length in the cryptanalysis of the Vigenère cipher.

The `calculate_matches` kernel is designed to calculate the number of matching characters for each possible shift value across the cipher text. Each thread handles a different shift value, iterating through the cipher text and incrementing a counter whenever a character matches its shifted counterpart.

4)Generating Key(crack_vignere_kernel):

Once the length of the key is known, the `crack_vignere_cuda` function and `crack_vignere_kernel` break the encrypted text into n Caesar Ciphers and Threads iterate through their assigned sections, counting the frequency of each letter. The frequency analysis aims to identify the most common letter in each segment, under the assumption that it represents the encrypted form of the letter 'E'—the most frequently occurring letter in English text.

Optimization Strategies

1)Minimizing Global Memory Access Latency: The kernel accesses the cipher text stored in global memory. To reduce the latency, the text is accessed sequentially by threads, which is beneficial for memory coalescing, where the GPU can efficiently load contiguous memory segments into the cache, reducing the number of required memory transactions.

2)Balancing Load Across Threads: The workload is evenly distributed among threads to ensure all are actively computing matches for different shift values, minimizing idle time and maximizing hardware utilization.

3)Using Atomic Operations for Counting Matches: To avoid race conditions and ensure correct match counts in parallel executions, atomic operations could be utilized when incrementing the match counters in global memory.

Memory Management:

1)Wherever Possible use global memory and `cudaMemcpyToSymbol` when the same resources need to be accessed by all threads, especially for data that needs to be accessed by all threads, such as the encryption key during encryption and decryption processes. This approach ensures efficient use of the GPU's memory hierarchy.

2) **Resource Cleanup:** Free the Cuda Memory when the job is done

Execution Time Measurements:

CPU-based Implementation (Standard C): The Normal C code takes approximately 0.189364 seconds for a sample c code.

CUDA-accelerated Implementation: Utilizing CUDA to parallelize the cipher operations significantly reduces the execution time to 0.122261 seconds.

It can be observed that CUDA implementation can exhibits a 35.44% improvement in execution time over the CPU-based implementation..

Bottlenecks and Limitations:

- 1) **GPU Utilization:** For smaller workloads, the overhead of initializing CUDA kernels and managing GPU resources may not be fully offset by the gains in parallel execution, leading to smaller speedups.
- 2) **Limitation of Kaiski-algorithm :** The project depends on finding the Kaiski-algorithm for finding the length of the key but as the length of the key increases the efficiency of the Kaiski-algorithm decreases. So all the limitations of Kaiski algorithm are also inherited in this code.

References:

- 1) Impact of Computational Power on Cryptography by Bhat, M.I., Giri, K.J. (2021)
DOI: 10.1007/978-981-15-8711-5_4
- 2) Analyzing the Kasiski Method Against Vigenere Cipher by April Lia Hananto
DOI: 10.48550/arXiv.1912.04519
- 3) <https://youtu.be/QgHnr8-h0xI>

Conclusion

The CUDA-based implementation for encrypting, decrypting, and cracking Vigenère cipher demonstrates a very efficient than classical C program increasing the efficiency by 35% . By leveraging the computational power of GPUs through CUDA programming, the project significantly enhances the speed and efficiency of the cryptographic processes involved.